

# Technical Handbook

Version 1.0 · 2025

---

|                 |                                                             |
|-----------------|-------------------------------------------------------------|
| <b>Audience</b> | Administrators, Developers, Project Managers                |
| <b>System</b>   | Infrastructure Webshop                                      |
| <b>Stack</b>    | Go · PostgreSQL · HTMX · OpenTofu · GitLab · Linode         |
| <b>Scope</b>    | Architecture · Operations · Development · Product Lifecycle |

---

# Contents

|          |                                                         |           |
|----------|---------------------------------------------------------|-----------|
| <b>I</b> | <b>Overview &amp; Architecture</b>                      | <b>9</b>  |
| <b>1</b> | <b>Introduction</b>                                     | <b>10</b> |
| 1.1      | Purpose of This Handbook . . . . .                      | 10        |
| 1.2      | Target Audiences . . . . .                              | 10        |
| 1.3      | Conventions Used in This Handbook . . . . .             | 11        |
| 1.4      | System At a Glance . . . . .                            | 11        |
| <b>2</b> | <b>C4 Architecture Model</b>                            | <b>12</b> |
| 2.1      | Level 1 — System Context . . . . .                      | 12        |
| 2.1.1    | Actors . . . . .                                        | 13        |
| 2.1.2    | External Systems . . . . .                              | 13        |
| 2.2      | Level 2 — Container Diagram . . . . .                   | 14        |
| 2.3      | Level 3 — Component Diagram . . . . .                   | 14        |
| 2.3.1    | Component Descriptions . . . . .                        | 15        |
| 2.4      | Level 4 — Key Code Flows . . . . .                      | 16        |
| 2.4.1    | Order Placement Flow (Project Manager) . . . . .        | 16        |
| 2.4.2    | Order Approval & Provisioning Flow . . . . .            | 17        |
| 2.4.3    | Request-Response Lifecycle (HTMX) . . . . .             | 17        |
| <b>3</b> | <b>Design Decisions &amp; Technology Rationale</b>      | <b>18</b> |
| 3.1      | Go as Application Language . . . . .                    | 18        |
| 3.2      | Server-Side Rendering with HTMX . . . . .               | 18        |
| 3.3      | DaisyUI + Tailwind CSS . . . . .                        | 19        |
| 3.4      | PostgreSQL with pgx (No ORM) . . . . .                  | 19        |
| 3.5      | Single-Container Architecture . . . . .                 | 19        |
| 3.6      | OpenTofu over Terraform . . . . .                       | 20        |
| 3.7      | GitLab Managed Terraform State . . . . .                | 20        |
| 3.8      | Per-Resource Isolated State . . . . .                   | 20        |
| 3.9      | Webhook-Driven Pipeline Trigger . . . . .               | 20        |
| 3.10     | Internationalization: 25 Languages . . . . .            | 21        |
| 3.11     | Session Security: Encrypted Cookies . . . . .           | 21        |
| 3.12     | Password Hashing: bcrypt . . . . .                      | 21        |
| 3.13     | Static Asset Embedding: <code>embed.FS</code> . . . . . | 22        |

|            |                                                |           |
|------------|------------------------------------------------|-----------|
| <b>II</b>  | <b>Infrastructure as Code</b>                  | <b>23</b> |
| <b>4</b>   | <b>GitLab Repository Concept</b>               | <b>24</b> |
| 4.1        | Overview . . . . .                             | 24        |
| 4.2        | Repository File Structure . . . . .            | 24        |
| 4.3        | State Management . . . . .                     | 25        |
| 4.4        | CI/CD Pipeline Stages . . . . .                | 25        |
| 4.5        | Webhook Integration with the Webshop . . . . . | 28        |
| 4.6        | Dependency Between Modules . . . . .           | 29        |
| <b>5</b>   | <b>OpenTofu Modules</b>                        | <b>30</b> |
| 5.1        | network-core . . . . .                         | 30        |
| 5.2        | vm-platform . . . . .                          | 30        |
| 5.3        | k8s-platform . . . . .                         | 31        |
| 5.4        | shared-services . . . . .                      | 31        |
| 5.5        | observability . . . . .                        | 31        |
| <b>III</b> | <b>Administrator Guide</b>                     | <b>32</b> |
| <b>6</b>   | <b>Initial Setup &amp; Installation</b>        | <b>33</b> |
| 6.1        | Prerequisites . . . . .                        | 33        |
| 6.2        | Environment Variables . . . . .                | 33        |
| 6.3        | Deployment on a Docker Host . . . . .          | 34        |
| 6.4        | Deployment on Kubernetes . . . . .             | 37        |
| 6.5        | Database Migrations . . . . .                  | 42        |
| 6.6        | First Login . . . . .                          | 43        |
| <b>7</b>   | <b>System Configuration</b>                    | <b>44</b> |
| 7.1        | GitLab Sources . . . . .                       | 44        |
| 7.2        | Deployment Environments . . . . .              | 44        |
| 7.3        | SMTP Configuration . . . . .                   | 44        |
| 7.4        | AI Translation . . . . .                       | 45        |
| 7.5        | Branding . . . . .                             | 45        |
| 7.6        | Currencies . . . . .                           | 45        |
| <b>8</b>   | <b>User Management</b>                         | <b>46</b> |
| 8.1        | Roles . . . . .                                | 46        |
| 8.2        | Creating a Local Account . . . . .             | 46        |
| 8.3        | Editing Users . . . . .                        | 46        |
| 8.4        | Deactivating Users . . . . .                   | 47        |
| <b>9</b>   | <b>Audit Log</b>                               | <b>48</b> |
| 9.1        | Logged Events . . . . .                        | 48        |
| 9.2        | Filtering . . . . .                            | 48        |
| 9.3        | Export . . . . .                               | 49        |

|           |                                                      |           |
|-----------|------------------------------------------------------|-----------|
| <b>IV</b> | <b>Developer Guide</b>                               | <b>50</b> |
| <b>10</b> | <b>Development Environment</b>                       | <b>51</b> |
| 10.1      | Prerequisites . . . . .                              | 51        |
| 10.2      | Quick Start . . . . .                                | 51        |
| 10.3      | Makefile Targets . . . . .                           | 52        |
| <b>11</b> | <b>Codebase Structure</b>                            | <b>53</b> |
| 11.1      | Directory Layout . . . . .                           | 53        |
| 11.2      | Layered Architecture . . . . .                       | 53        |
| 11.3      | Template System . . . . .                            | 54        |
| 11.4      | Adding Translations . . . . .                        | 55        |
| 11.5      | Database Migrations . . . . .                        | 55        |
| <b>12</b> | <b>Adding a New Feature</b>                          | <b>56</b> |
| 12.1      | Example: Adding a New Admin Page . . . . .           | 56        |
| 12.1.1    | Step 1: Create the Migration . . . . .               | 56        |
| 12.1.2    | Step 2: Add the Model . . . . .                      | 56        |
| 12.1.3    | Step 3: Define the Repository Interface . . . . .    | 56        |
| 12.1.4    | Step 4: Implement the Repository . . . . .           | 57        |
| 12.1.5    | Step 5: Add the Handler . . . . .                    | 57        |
| 12.1.6    | Step 6: Create the Template . . . . .                | 57        |
| 12.1.7    | Step 7: Add i18n Keys . . . . .                      | 58        |
| 12.1.8    | Step 8: Wire Up in main.go . . . . .                 | 58        |
| 12.1.9    | Step 9: Test . . . . .                               | 58        |
| <b>V</b>  | <b>Project Manager &amp; Product Lifecycle Guide</b> | <b>59</b> |
| <b>13</b> | <b>Ordering Infrastructure</b>                       | <b>60</b> |
| 13.1      | The Order Lifecycle . . . . .                        | 60        |
| 13.2      | Placing an Order . . . . .                           | 60        |
| 13.3      | Tracking Status . . . . .                            | 61        |
| 13.4      | Using an Existing Deployment as a Template . . . . . | 61        |
| 13.5      | Decommissioning . . . . .                            | 61        |
| <b>14</b> | <b>Introducing a New Product</b>                     | <b>62</b> |
| 14.1      | Overview . . . . .                                   | 62        |
| 14.2      | Step 1: Plan the Product . . . . .                   | 62        |
| 14.3      | Step 2: Prepare the OpenTofu Module . . . . .        | 63        |
| 14.4      | Step 3: Configure the GitLab Webhook . . . . .       | 63        |
| 14.5      | Step 4: Create a Category (if needed) . . . . .      | 64        |
| 14.6      | Step 5: Create the Product . . . . .                 | 64        |
| 14.6.1    | Basic Information . . . . .                          | 64        |
| 14.6.2    | Parameters . . . . .                                 | 64        |
| 14.6.3    | Deployment Environments . . . . .                    | 65        |
| 14.6.4    | Multiple Webhooks (Pipeline Arrays) . . . . .        | 65        |

|                                                             |           |
|-------------------------------------------------------------|-----------|
| 14.7 Step 6: Configure AI Translation . . . . .             | 65        |
| 14.8 Step 7: Global and Category Parameter Sets . . . . .   | 66        |
| 14.9 Step 8: Test the Full Order Flow . . . . .             | 66        |
| 14.10 Step 9: Test with a Project Manager Account . . . . . | 67        |
| 14.11 Step 10: Product is Live . . . . .                    | 67        |
| <b>15 Project Management</b>                                | <b>68</b> |
| 15.1 Projects . . . . .                                     | 68        |
| 15.2 Infrastructure Overview . . . . .                      | 68        |
| 15.3 Cost Center Assignment . . . . .                       | 68        |
| <b>A Supported Languages</b>                                | <b>70</b> |
| <b>B Environment Variable Reference</b>                     | <b>71</b> |
| <b>C OpenTofu Module Variable Reference</b>                 | <b>72</b> |
| C.1 network-core Variables . . . . .                        | 72        |
| C.2 vm-platform Variables . . . . .                         | 72        |
| C.3 k8s-platform Variables . . . . .                        | 73        |
| C.4 shared-services Variables . . . . .                     | 73        |
| C.5 observability Variables . . . . .                       | 74        |
| <b>D HTTP Route Reference</b>                               | <b>75</b> |
| <b>E Requirements Traceability Matrix</b>                   | <b>77</b> |
| <b>F Compiling This Handbook</b>                            | <b>79</b> |

# List of Figures

|      |                                                                |    |
|------|----------------------------------------------------------------|----|
| 2.1  | C4 Level 1 — System Context Diagram . . . . .                  | 12 |
| 2.2  | C4 Level 2 — Container Diagram . . . . .                       | 14 |
| 2.3  | C4 Level 3 — Component Diagram (Webshop Application) . . . . . | 15 |
| 4.1  | Module dependency graph . . . . .                              | 29 |
| 6.1  | Docker Host deployment architecture . . . . .                  | 35 |
| 6.2  | Kubernetes deployment architecture . . . . .                   | 37 |
| 13.1 | Order lifecycle state machine . . . . .                        | 60 |

# List of Tables

|      |                                            |    |
|------|--------------------------------------------|----|
| 1.1  | Technology stack overview . . . . .        | 11 |
| 2.1  | Component descriptions . . . . .           | 16 |
| 4.1  | GitLab repository overview . . . . .       | 24 |
| 4.2  | State name convention per module . . . . . | 25 |
| 6.1  | Required environment variables . . . . .   | 34 |
| 6.2  | Key Helm values . . . . .                  | 39 |
| 8.1  | User roles and capabilities . . . . .      | 46 |
| 9.1  | Audited actions . . . . .                  | 48 |
| 10.1 | Makefile targets . . . . .                 | 52 |
| A.1  | Supported language codes . . . . .         | 70 |
| B.1  | Environment variable reference . . . . .   | 71 |
| C.1  | network-core variables . . . . .           | 72 |
| C.2  | vm-platform variables . . . . .            | 73 |
| C.3  | k8s-platform variables . . . . .           | 73 |
| C.4  | shared-services variables . . . . .        | 74 |
| C.5  | observability variables . . . . .          | 74 |
| D.1  | HTTP route reference (selected) . . . . .  | 76 |
| E.1  | Requirements traceability matrix . . . . . | 78 |

# Listings

|     |                                                           |    |
|-----|-----------------------------------------------------------|----|
| 4.1 | .gitlab-ci.yml for vm-platform (webhook-triggered module) | 25 |
| 4.2 | .gitlab-ci.yml for singleton modules (manual apply)       | 27 |
| 6.1 | deploy/docker-host/nginx.conf                             | 36 |
| 6.2 | Helm install with required values                         | 38 |
| 6.3 | Example secrets.values.yaml (do not commit)               | 38 |
| 6.4 | Kubernetes Secret for the webshop                         | 39 |
| 6.5 | Migration Job                                             | 40 |
| 6.6 | Deployment and Service                                    | 40 |
| 6.7 | Ingress with cert-manager TLS                             | 41 |
| 6.8 | HorizontalPodAutoscaler                                   | 42 |

# Part I

## Overview & Architecture

# Chapter 1

## Introduction

### 1.1 Purpose of This Handbook

The PORR AG Infrastructure Webshop is a self-service portal that enables Digital Unit (DU) Admins and Project Managers to order, manage, and decommission IT infrastructure on Linode/Akamai Cloud through a browser-based interface. This handbook is the authoritative reference for everyone who operates, extends, or uses the system.

It covers:

- The complete system architecture described with the C4 model
- Infrastructure-as-Code modules and the GitLab CI/CD concept
- Architecture and design decisions with their rationale
- Step-by-step administrator operations
- Developer onboarding, code structure, and extension patterns
- The full product lifecycle from creation to decommissioning

### 1.2 Target Audiences

|                        |                                                                                                                                                                              |
|------------------------|------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>DU Admin</b>        | Responsible for approving orders, monitoring all infrastructure, and ordering directly without an approval step. Reads Part <a href="#">III</a> and Part <a href="#">V</a> . |
| <b>Webshop Admin</b>   | Responsible for the product catalog, system configuration, and local user accounts. Reads Part <a href="#">III</a> entirely.                                                 |
| <b>Developer</b>       | Extends or maintains the webshop codebase and the OpenTofu modules. Reads Part <a href="#">I</a> and Part <a href="#">IV</a> .                                               |
| <b>Project Manager</b> | Places infrastructure orders and manages projects. Reads Part <a href="#">V</a> .                                                                                            |

## 1.3 Conventions Used in This Handbook

**monospace** File paths, commands, code, environment variables, and route patterns.

***italic*** Technical terms on first use.

**bold** UI labels, button names, and menu items.

**Note boxes** Yellow-highlighted boxes providing additional context or important background information.

**Important boxes** Orange-highlighted boxes warning about actions that can cause data loss or service disruption.

**Tip boxes** Green-highlighted boxes with shortcuts and best practices.

## 1.4 System At a Glance

The webshop is a single Go binary that renders HTML server-side and delivers page fragments via HTMX for a responsive, app-like feel. Orders are fulfilled by triggering GitLab CI/CD pipelines that run OpenTofu (a Terraform-compatible open-source tool) to provision or destroy cloud resources on Linode/Akamai. The system is stateless above the PostgreSQL layer, enabling horizontal scaling without coordination.

Table 1.1: Technology stack overview

| Layer         | Technology              | Version    |
|---------------|-------------------------|------------|
| Backend       | Go                      | 1.22+      |
| Frontend      | HTMX + Alpine.js        | 2.0 / 3.14 |
| CSS           | Tailwind CSS + DaisyUI  | 3.x / 4.x  |
| Database      | PostgreSQL              | 15+        |
| DB Driver     | pgx/v5 (jackc)          | 5.x        |
| IaC           | OpenTofu                | 1.9+       |
| Cloud         | Linode/Akamai           | —          |
| CI/CD         | GitLab CE               | 16+        |
| State Backend | GitLab Managed TF State | —          |
| Auth          | OIDC (Entra ID) + local | —          |

# Chapter 2

## C4 Architecture Model

The system is described using the *C4 model* (Context, Container, Component, Code). Each level zooms in one step deeper.

### 2.1 Level 1 — System Context

Figure 2.1 shows who uses the webshop and which external systems it depends on.




Figure 2.1: C4 Level 1 — System Context Diagram

### 2.1.1 Actors

**DU Admin (Digital Unit Administrator)** Authenticates via Microsoft Entra ID SSO. Has the system role `admin`. Can order infrastructure directly (no approval required), approve or reject orders from Project Managers, view all projects and infrastructure, and decommission any resource.

**Webshop Admin** Uses a local account (no SSO). Has the system role `shop_admin`. Manages the product catalog, system configuration (SMTP, AI, branding, currencies), environments, GitLab sources, cost centers, and local users. Cannot place orders.

**Project Manager** Authenticates via Entra ID SSO. Has the system role `user`. Places orders (which require DU Admin approval), manages their own projects, and can decommission their own infrastructure.

### 2.1.2 External Systems

**GitLab** One or more self-hosted or cloud GitLab instances. The webshop connects to GitLab for three purposes: (1) browsing repositories and importing `variables.tf` parameter definitions, (2) triggering provisioning and decommissioning webhooks, and (3) polling pipeline status to update order state.

**Microsoft Entra ID** The OIDC identity provider for DU Admins and Project Managers. The webshop implements the Authorization Code Flow. Local accounts for Webshop Admins bypass this provider entirely.

**AI Translation API** Optional. Used to auto-translate product names and descriptions into all 25 supported languages. Multiple backends are supported: cloud-based (Claude, OpenAI, Azure OpenAI) and on-premise (Ollama, LocalAI). The feature is hidden when no provider is configured.

**Mail Server** Any SMTP server. Sends order confirmations, approval request notifications, approval/rejection notices, deployment completion alerts, and failure notifications.

## 2.2 Level 2 — Container Diagram




Figure 2.2: C4 Level 2 — Container Diagram

The system has exactly two containers:

**Webshop Application** A single stateless Go binary. It handles HTTP requests, renders HTML via `html/template`, serves HTMX fragments, coordinates all business logic, runs the GitLab polling goroutines, and manages outbound SMTP. Being stateless above the database layer means multiple replicas can run behind a load balancer without any coordination layer.

**PostgreSQL Database** The single source of truth. Stores the product catalog, orders, infrastructure state, audit log, branding, configuration, users, and exchange rates. The database is the only stateful component.

## 2.3 Level 3 — Component Diagram

Figure 2.3 decomposes the Go application container into its major logical components.




Figure 2.3: C4 Level 3 — Component Diagram (Webshop Application)

### 2.3.1 Component Descriptions

| Component      | Responsibility                                                                                                                                                                                                                                  |
|----------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Auth           | Handles OIDC Authorization Code Flow (Entra ID) and local password login. Issues encrypted, HttpOnly session cookies. Enforces role-based access ( <code>admin</code> , <code>shop_admin</code> , <code>user</code> ).                          |
| Catalog        | Renders products grouped by category. Converts prices from the configured base currency to the user's locale currency using stored exchange rates. Supports 25 languages via the <code>i18n</code> engine.                                      |
| Order          | Renders the order form with parameters from global, category, and product parameter sets. Stores orders in PostgreSQL. For DU Admins, triggers the GitLab provisioning webhook immediately; for Project Managers, parks the order for approval. |
| Approval       | Lists pending orders for DU Admins. Approve triggers the GitLab webhook. Reject requires a mandatory comment. Both paths send email notifications and write to the audit log.                                                                   |
| Infrastructure | Shows deployed resources grouped by project and environment. Role-scoped: Project Managers see only their own. Triggers the GitLab destroy webhook for decommissioning. Supports “use as template” to pre-fill a new order.                     |

| Component      | Responsibility                                                                                                                                                                                                                                                      |
|----------------|---------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Audit          | Append-only log of all system actions. Filterable by date range, user, and action type. Exportable as CSV or PDF.                                                                                                                                                   |
| Notification   | Wraps SMTP email delivery. Silently no-ops when SMTP is not configured. SMTP credentials are live-reloadable from the database without a server restart.                                                                                                            |
| GitLab Polling | Goroutine pool that periodically queries the GitLab Pipelines API for in-flight orders. Updates order and infrastructure state on completion or failure, then triggers email notifications. Uses PostgreSQL advisory locks so multiple replicas do not double-poll. |
| Admin          | All Webshop Admin pages: categories, products (CRUD, image upload, AI translation, webhook management, parameter import from <code>variables.tf</code> ), environments, GitLab sources, cost centers, currencies, SMTP, AI config, branding, users.                 |
| AI Translation | Pluggable interface ( <code>aitranslation.Translator</code> ). Concrete implementations for Claude, OpenAI, Azure OpenAI, Ollama, and LocalAI. Selected and instantiated at startup from the database configuration; swappable at runtime via the Admin UI.         |
| Repositories   | Thin SQL layer using <code>pgx/v5</code> . One repository interface per aggregate root (User, Product, Order, etc.). Implementations in <a href="#">internal/repository/postgres/</a> .                                                                             |

Table 2.1: Component descriptions

## 2.4 Level 4 — Key Code Flows

### 2.4.1 Order Placement Flow (Project Manager)

1. Browser submits `POST /orders/new`.
2. `handler.orderCreate()` validates input, calls `service.OrderService.Create()`.
3. The order service writes the order record with status `pending_approval` to PostgreSQL.
4. The notification service sends an approval-request email to all DU Admins.
5. The audit service writes an `order_placed` entry.
6. Browser redirects to `GET /orders/{id}` showing the *Awaiting Approval* status.

### 2.4.2 Order Approval & Provisioning Flow

1. DU Admin opens `GET /approvals` and clicks **Approve**.
2. `handler.approvalApprove()` calls `service.OrderService.Approve()`.
3. The order service updates status to `provisioning`, stores the webhook URLs from the product-environment configuration.
4. For each configured webhook (ordered by `exec_order`): an HTTP POST is sent to the GitLab webhook URL with the Linode/Tofu parameters as JSON.
5. GitLab receives the trigger, starts a pipeline, returns a pipeline ID.
6. The webshop stores the pipeline ID(s) in the JSONB `pipeline_ids` column.
7. The GitLab Polling goroutine detects the in-flight pipeline and starts querying the Pipelines API every 30 seconds.
8. When the pipeline completes successfully, the poller sets status to `completed` and sends a deployment-complete email.
9. On failure, status becomes `failed` and failure emails go to the orderer and all DU Admins.

### 2.4.3 Request–Response Lifecycle (HTMX)

Most user interactions either perform a full page load (navigation) or return an HTML fragment (HTMX). The server detects HTMX requests via the `HX-Request` header. Full renders use the layout template with content injected into the `{{block "content"}}` slot. Fragments skip the layout and return only the relevant HTML snippet, which HTMX swaps into the DOM.

# Chapter 3

## Design Decisions & Technology Rationale

### 3.1 Go as Application Language

**Decision:** The entire backend is written in Go.

**Rationale:**

- Compiles to a single self-contained binary with no runtime dependencies — ideal for Docker-based deployment.
- Low memory footprint and predictable garbage collection enable running multiple polling goroutines without resource pressure.
- The standard library `net/http` in Go 1.22+ provides pattern-matching routing (`"GET /path/{id}"`) sufficient for this use case, removing the need for a router dependency.
- `html/template` provides automatic context-aware escaping, eliminating entire classes of XSS vulnerabilities.
- Goroutines provide a natural model for the concurrent GitLab pipeline polling without callback complexity.

**Alternative considered:** Python/FastAPI with Jinja2. Rejected because the deployment story is more complex (virtualenv, interpreter) and concurrency requires `async`, which increases cognitive overhead.

### 3.2 Server-Side Rendering with HTMX

**Decision:** Pages are rendered server-side; dynamic updates use HTMX HTML fragments rather than a JavaScript SPA framework.

**Rationale:**

- No JSON API surface to design, document, or secure separately — the server owns both data and presentation.
- Progressive enhancement: the UI degrades gracefully if JS is unavailable.
- HTMX's `hx-get/hx-post` attributes on HTML elements cover the interactivity requirements (live status polling, inline form submission, partial page swaps) without React or Vue build pipelines.
- Template-level security: Go's `html/template` auto-escapes all inserted values, preventing XSS by construction.

### 3.3 DaisyUI + Tailwind CSS

**Decision:** Tailwind CSS for utility classes, DaisyUI for semantic component classes (`btn`, `badge`, `card`, `modal`, etc.).

**Rationale:** DaisyUI components map directly to HTML elements and class names — no JavaScript component framework is required. The combination produces accessible, consistent UI with minimal custom CSS. Tailwind's JIT compilation trims the final CSS file to only used classes.

### 3.4 PostgreSQL with pgx (No ORM)

**Decision:** PostgreSQL 15+, accessed via `pgx/v5` with hand-written SQL. No ORM.

**Rationale:**

- JSONB arrays for `pipeline_ids` allow storing multiple concurrent pipeline references per order without a join table, while retaining full query ability.
- BYTEA columns store product images and the branding logo directly in the database, eliminating a separate object store for this scale.
- Direct SQL gives precise control over query shape and enables `ON CONFLICT DO UPDATE` (upserts) for the branding and app-config singleton rows.
- PostgreSQL advisory locks coordinate the polling goroutines across multiple replicas without an additional message broker.
- `pgx` is the highest-performance PostgreSQL driver for Go; it avoids the `database/sql` reflection overhead.

### 3.5 Single-Container Architecture

**Decision:** The entire application (HTTP server, background polling, notification dispatch) runs in a single container.

**Rationale:** At the anticipated load (tens of users, hundreds of orders per month), microservice decomposition would add operational complexity (service discovery, inter-service auth, distributed tracing) without meaningful benefit. The single-binary model keeps deployment simple: one `docker run` or one Kubernetes `Deployment` object. Horizontal scaling is achieved by running multiple replicas; PostgreSQL advisory locks ensure the polling goroutines do not duplicate work.

## 3.6 OpenTofu over Terraform

**Decision:** All IaC modules use OpenTofu 1.9+.

**Rationale:** OpenTofu is a community-governed, truly open-source fork of Terraform created after HashiCorp changed Terraform’s license to BSL 1.1 in 2023. The API, HCL syntax, provider ecosystem, and state format are fully compatible. Using OpenTofu avoids future licensing costs and complies with PORR’s open-source software policy.

## 3.7 GitLab Managed Terraform State

**Decision:** OpenTofu state is stored in GitLab’s built-in HTTP-compatible state backend (GitLab Managed Terraform State).

**Rationale:**

- Reuses existing GitLab infrastructure — no S3 bucket, no Terraform Cloud account, no Vault setup required.
- State is protected by GitLab’s RBAC: only project members with Developer or higher access can read or modify state.
- State locking via HTTP `LOCK/UNLOCK` prevents concurrent apply operations.
- State is browsable in the GitLab UI under **Infrastructure** → **Terraform states**.

## 3.8 Per-Resource Isolated State

**Decision:** Each VM and each Kubernetes cluster has its own isolated OpenTofu state (named `vm-{label}` and `k8s-{label}`).

**Rationale:** Shared state across multiple VMs creates blast-radius risk: a failing `apply` for one VM could lock the state and block unrelated VMs. Isolated states mean each resource’s lifecycle is entirely independent. The trade-off is a larger number of state files, which is manageable at the expected scale.

## 3.9 Webhook-Driven Pipeline Trigger

**Decision:** The webshop triggers GitLab pipelines via webhook (HTTP POST) rather than the GitLab API’s pipeline-trigger endpoint.

**Rationale:** Webhooks are the native GitLab mechanism for external triggers. They carry a secure token for authentication and include the ability to pass variables (OpenTofu parameters) as JSON. The pipeline configuration ([.gitlab-ci.yml](#)) lives in the same repository as the IaC code, making the trigger configuration auditable and version-controlled.

## 3.10 Internationalization: 25 Languages

**Decision:** All UI strings are externalized in a key-value i18n map covering 25 languages (all EU official languages plus Russian and Ukrainian).

**Rationale:** PORR AG operates across Europe. The map is defined in [internal/i18n/i18n.go](#) as a Go map literal for compile-time completeness checking. Translation of product content (not UI strings) is delegated to the optional AI translation service.

## 3.11 Session Security: Encrypted Cookies

**Decision:** User sessions are stored in encrypted, signed HTTP cookies — not in the database or server-side session store.

**Rationale:**

- No database round-trip on every request to look up a session.
- Horizontal scaling is trivial: no shared session store required.
- The `SESSION_KEY` (32-byte, hex-encoded) is used to both sign and encrypt the cookie content via AES-GCM, so the server can verify integrity without a separate HMAC key.
- Cookies are set with `HttpOnly`, `Secure`, and `SameSite=Lax` to mitigate XSS session theft and CSRF.
- Logout invalidates the session by overwriting the cookie with an empty, expired value — no server-side revocation list needed for the MVP threat model.

**Alternative considered:** Server-side sessions in PostgreSQL. Rejected to avoid the extra table and join on every authenticated request. If revocation of individual sessions at scale is required this can be introduced as a future enhancement.

## 3.12 Password Hashing: bcrypt

**Decision:** Local user passwords are hashed with bcrypt (cost factor 12) via [golang.org/x/crypto/bcrypt](#).

**Rationale:**

- bcrypt is an adaptive, memory-hard function specifically designed for password storage; it is resistant to GPU-accelerated dictionary attacks.

- Cost factor 12 takes 250 ms on modern hardware — slow enough to deter brute-force while still acceptable for interactive login.
- The Go standard-library-adjacent `x/crypto` package is actively maintained by the Go team.
- Passwords are never stored in plaintext or with reversible encryption.

**Alternative considered:** Argon2id. More modern, but `bcrypt` is widely understood, well-tested in Go, and sufficient for the low login-volume use case.

### 3.13 Static Asset Embedding: `embed.FS`

**Decision:** HTML templates and static assets (CSS, JS, images) are embedded into the Go binary at compile time using the `//go:embed` directive and `embed.FS`.

**Rationale:**

- The resulting binary is fully self-contained — no external file system paths need to match the deployment environment.
- Deployment is a single `COPY` in the Dockerfile (multi-stage build); no volume mounts are required for templates.
- Template changes are compiled into the binary, so a deployment is always atomic: the binary and its templates are in sync.
- The Go compiler resolves embed paths at build time, making missing-file bugs a compile error rather than a runtime panic.

**Trade-off:** Hot-reloading templates during development is not possible without restarting the server. The `make dev` target runs the server with `go run`, so a simple restart is fast enough for the development workflow.

## Part II

# Infrastructure as Code

# Chapter 4

## GitLab Repository Concept

### 4.1 Overview

The IaC layer consists of six GitLab repositories in the `porr-ag` group. Each repository is a self-contained OpenTofu module with its own CI/CD pipeline and state. Together they provision all cloud infrastructure on Linode/Akamai.

Table 4.1: GitLab repository overview

| Repository                   | Purpose                               | Trigger |
|------------------------------|---------------------------------------|---------|
| <code>network-core</code>    | VPC, subnets, base firewall, DNS zone | Manual  |
| <code>vm-platform</code>     | Linode VM instances (per-VM)          | Webhook |
| <code>k8s-platform</code>    | LKE Kubernetes clusters (per-cluster) | Webhook |
| <code>shared-services</code> | GitLab CE + Vault + Registry          | Manual  |
| <code>observability</code>   | Grafana + Prometheus + Loki           | Manual  |
| <code>infra-state</code>     | State backend host (no IaC code)      | N/A     |

### 4.2 Repository File Structure

Every module repository (except `infra-state`) follows the same file layout:

```
<module>/
+-- .gitlab-ci.yml    # CI/CD pipeline definition
+-- backend.tf        # HTTP state backend (GitLab Managed)
+-- provider.tf       # Linode provider + required versions
+-- variables.tf      # Input variable declarations
+-- main.tf           # Resource definitions
+-- outputs.tf        # Output values
```

## 4.3 State Management

All state files live in the `infra-state` project. The project hosts no IaC code; it is solely a container for GitLab Managed Terraform State files.

**State naming convention:**

Table 4.2: State name convention per module

| Module                       | State name                       |
|------------------------------|----------------------------------|
| <code>network-core</code>    | <code>network-core</code>        |
| <code>vm-platform</code>     | <code>vm-{vm_label}</code>       |
| <code>k8s-platform</code>    | <code>k8s-{cluster_label}</code> |
| <code>shared-services</code> | <code>shared-prod</code>         |
| <code>observability</code>   | <code>observability-prod</code>  |

The backend URL pattern is:

```
https://<gitlab-host>/api/v4/projects/<INFRA_STATE_PROJECT_ID>/terraform/state/<
```

Backend configuration is passed at `tofu init` time via `-backend-config` flags, not hard-coded in `backend.tf`. This keeps credentials out of version control.

## 4.4 CI/CD Pipeline Stages

All module pipelines share the same three-stage structure:

1. **validate** — `tofu validate`: checks HCL syntax and provider schema compliance.
2. **plan** — `tofu plan -out=tfplan`: computes the diff and saves the plan artifact.
3. **apply** — `tofu apply tfplan`: applies the plan. For webhook-triggered modules, apply runs automatically on pipeline success. For singleton modules (`shared-services`, `observability`, `network-core`), the apply job is `when: manual` to require deliberate human confirmation.

**Destroy flow:** Webhook-triggered modules (`vm-platform`, `k8s-platform`) accept a `TF_ACTION=destroy` variable. When set, the apply stage runs `tofu destroy -auto-approve` instead of `tofu apply`.

Listing 4.1 shows the canonical `.gitlab-ci.yml` for the `vm-platform` module. Singleton modules differ only in the apply job: `when: manual` replaces `when: on_success`, and there is no destroy branch.

```
image: ghcr.io/opentofu/opentofu:1.9

variables:
  TF_ACTION: "apply"
  TF_STATE_NAME: "vm-${TF_VAR_vm_label}"
```

```
TF_VAR_linode_token: "${LINODE_TOKEN}"

.tofu_init: &tofu_init
  before_script:
    - tofu init
      -reconfigure
      -backend-config="address=${CI_API_V4_URL}/projects/${INFRA_STATE_PROJECT_ID}"
      -backend-config="lock_address=${CI_API_V4_URL}/projects/${INFRA_STATE_PROJECT_ID}"
      -backend-config="unlock_address=${CI_API_V4_URL}/projects/${INFRA_STATE_PROJECT_ID}"
      -backend-config="username=gitlab-ci-token"
      -backend-config="password=${CI_JOB_TOKEN}"
      -backend-config="lock_method=POST"
      -backend-config="unlock_method=DELETE"
      -backend-config="retry_wait_min=5"

  stages:
    - validate
    - plan
    - apply

  validate:
    stage: validate
    <<: *tofu_init
    script:
      - tofu validate
    rules:
      - if: $CI_PIPELINE_SOURCE == "trigger"

  plan:
    stage: plan
    <<: *tofu_init
    script:
      - tofu plan -out=tfplan
    artifacts:
      paths:
        - tfplan
      expire_in: 1 hour
    rules:
      - if: $CI_PIPELINE_SOURCE == "trigger"

  apply:
    stage: apply
    <<: *tofu_init
    script:
      - |
        if [ "${TF_ACTION}" = "destroy" ]; then
          tofu destroy -auto-approve
        else
          tofu apply -auto-approve tfplan
        fi
    dependencies:
      - plan
    rules:
      - if: $CI_PIPELINE_SOURCE == "trigger" && $TF_ACTION == "destroy"
      when: always
```

```
- if: $CI_PIPELINE_SOURCE == "trigger"
  when: on_success
```

Listing 4.1: .gitlab-ci.yml for vm-platform (webhook-triggered module)

Listing 4.2 shows a complete pipeline for singleton modules (`network-core`, `shared-services`, `observability`) where a human must confirm the apply step:

```
image: ghcr.io/opentofu/opentofu:1.9

variables:
  TF_STATE_NAME: "shared-prod"
  TF_VAR_linode_token: "${LINODE_TOKEN}"

.tofu_init: &tofu_init
  before_script:
    - tofu init
      - backend-config="address=${CI_API_V4_URL}/projects/${INFRA_STATE_PROJECT}"
      - backend-config="lock_address=${CI_API_V4_URL}/projects/${INFRA_STATE_PROJECT}"
      - backend-config="unlock_address=${CI_API_V4_URL}/projects/${INFRA_STATE_PROJECT}"
      - backend-config="username=gitlab-ci-token"
      - backend-config="password=${CI_JOB_TOKEN}"
      - backend-config="lock_method=POST"
      - backend-config="unlock_method=DELETE"
      - backend-config="retry_wait_min=5"

stages:
  - validate
  - plan
  - apply

validate:
  stage: validate
  <<: *tofu_init
  script:
    - tofu validate
  rules:
    - if: $CI_COMMIT_BRANCH == "main"

plan:
  stage: plan
  <<: *tofu_init
  script:
    - tofu plan -out=tfplan
  artifacts:
    paths:
      - tfplan
    expire_in: 1 hour
  rules:
    - if: $CI_COMMIT_BRANCH == "main"

apply:
  stage: apply
  <<: *tofu_init
  script:
    - tofu apply -auto-approve tfplan
```

```
dependencies:
  - plan
rules:
  - if: $CI_COMMIT_BRANCH == "main"
    when: manual
```

Listing 4.2: .gitlab-ci.yml for singleton modules (manual apply)

### Note

The single required CI/CD variable to set in the *infra-state* GitLab project group is `INFRA_STATE_PROJECT_ID` — the numeric project ID of the `infra-state` repository (visible in GitLab under **Settings** → **General**). `CI_JOB_TOKEN` and `CI_API_V4_URL` are provided automatically by GitLab CI. `LINODE_TOKEN` must be set with Read/Write API access.

## 4.5 Webhook Integration with the Webshop

1. A DU Admin approves an order in the webshop (or places one directly).
2. For each webhook configured on the product-environment combination, the webshop sends:

```
POST <gitlab-webhook-url>
X-Gitlab-Token: <secret>
Content-Type: application/json

{
  "ref": "main",
  "variables": [
    {"key": "TF_VAR_vm_label", "value": "myvm-01"},
    {"key": "TF_VAR_vpc_id", "value": "12345"},
    ...
  ]
}
```

3. GitLab starts a pipeline for the `main` branch, injecting the variables as environment variables.
4. The pipeline reads `TF_VAR_*` variables, runs `tofu plan` then `tofu apply`.
5. On completion (success or failure), the webshop's polling goroutine detects the status change and updates the order record.

## 4.6 Dependency Between Modules




Figure 4.1: Module dependency graph

**network-core** must be applied first. Its outputs (`vpc_id`, `subnet_ids`, `firewall_id`) are passed as input variables to all other modules.

# Chapter 5

## OpenTofu Modules

### 5.1 network-core

**Purpose:** Foundational network infrastructure. Apply once per environment; no webhook trigger.

**Resources created:** `linode_vpc` (VPC), `linode_vpc_subnet` (subnets via `for_each`), `linode_firewall` (base inbound/outbound policy), `linode_domain` (optional DNS zone).

**Key design choice:** Subnets are defined as a `map(object({cidr, description}))` variable. This allows adding subnets by editing a single variable without changing the resource code.

**Key outputs:** `vpc_id`, `subnet_ids` (map: label  $\rightarrow$  ID), `firewall_id`, `dns_zone_id`.

### 5.2 vm-platform

**Purpose:** Provision individual Linode VMs on demand. One state per VM (`vm-{label}`). Webhook-triggered by the webshop.

**Resources created:** `linode_firewall` (VM-specific), `linode_instance`, `linode_firewall_device` (attaches firewall to VM), `linode_volume` (optional block storage), `linode_domain_record` (optional A record).

**Key design choice:** The VM gets its own firewall (in addition to the base firewall from `network-core`) to allow per-VM port rules without affecting the shared base policy.

**Destroy flow:** Triggered by the webshop decommission action. `TF_ACTION=destroy`  $\rightarrow$  `tofu destroy -auto-approve`.

## 5.3 k8s-platform

**Purpose:** Provision Linode Kubernetes Engine (LKE) clusters. One state per cluster (`k8s-{label}`). Webhook-triggered.

**Resources created:** `linode_firewall` (cluster-specific), `linode_lke_cluster` with dynamic `pool` blocks (supports multiple node pools with optional autoscaler), `linode_domain_record` (optional).

**Key output:** `kubeconfig` (base64-encoded, marked `sensitive`) — retrieve with `tofu output -raw kubeconfig | base64 -d`.

**Node pool design:** Node pools are a `list(object)` variable with an optional `autoscaler` sub-object, allowing any combination of fixed-size and auto-scaling pools in a single cluster.

## 5.4 shared-services

**Purpose:** Long-lived shared services: GitLab CE, HashiCorp Vault, container registry. Single state (`shared-prod`). Manual apply.

**Resources created:** Two `linode_instance` resources (GitLab, Vault), two `linode_volume` resources (data volumes), one `linode_firewall` with tailored inbound rules (HTTP/HTTPS/Git SSH for GitLab; Vault API 8200 restricted to RFC1918), two `linode_firewall_device` attachments, optional DNS A records.

**Operational note:** GitLab CE and Vault are installed via post-provisioning scripts (cloud-init or Ansible), not managed by OpenTofu. OpenTofu only provisions the infrastructure layer.

## 5.5 observability

**Purpose:** Grafana, Prometheus, and Loki on a single VM. Single state (`observability-prod`). Manual apply.

**Resources created:** `linode_instance`, `linode_volume` (200 GB default for metrics and logs), `linode_firewall` (Grafana port configurable, Prometheus/Loki restricted to RFC1918), optional DNS A record.

**Key output:** `grafana_url` — `http://{ip}:{port}`.

# Part III

## Administrator Guide

# Chapter 6

## Initial Setup & Installation

### 6.1 Prerequisites

- Docker 24+ (or a Kubernetes cluster)
- PostgreSQL 15+ (external or containerized)
- SMTP server credentials (optional, can be configured later)
- Microsoft Entra ID application registration for OIDC (optional for local-only operation)
- Linode API token with Read/Write scope
- GitLab instance with webhook support

### 6.2 Environment Variables

The webshop is configured entirely via environment variables. Copy `.env.example` to `.env` and populate the following:

Table 6.1: Required environment variables

| Variable                        | Description                            | Required |
|---------------------------------|----------------------------------------|----------|
| <code>DATABASE_URL</code>       | PostgreSQL DSN                         | Yes      |
| <code>SESSION_KEY</code>        | 32-byte hex session encryption key     | Yes      |
| <code>ADMIN_EMAIL</code>        | Initial Webshop Admin email            | Yes      |
| <code>ADMIN_PASSWORD</code>     | Initial Webshop Admin password         | Yes      |
| <code>BASE_URL</code>           | Public URL of the webshop              | Yes      |
| <code>OIDC_CLIENT_ID</code>     | Entra ID application client ID         | No       |
| <code>OIDC_CLIENT_SECRET</code> | Entra ID client secret                 | No       |
| <code>OIDC_ISSUER</code>        | Entra ID issuer URL                    | No       |
| <code>SMTP_HOST</code>          | SMTP host (overridable in Admin UI)    | No       |
| <code>SMTP_PORT</code>          | SMTP port (default 587)                | No       |
| <code>SMTP_FROM</code>          | Sender address                         | No       |
| <code>SMTP_USERNAME</code>      | SMTP username                          | No       |
| <code>SMTP_PASSWORD</code>      | SMTP password                          | No       |
| <code>AI_PROVIDER</code>        | AI translation provider (startup only) | No       |

**Note**

SMTP and AI translation can also be configured (and overridden) through the Admin UI at `/admin/smtp` and `/admin/ai-config`. Database settings take precedence over environment variables at runtime.

## 6.3 Deployment on a Docker Host

This section describes a production-grade deployment on a single Linux host using Docker and an Nginx reverse proxy with TLS termination.

## Architecture




Figure 6.1: Docker Host deployment architecture

## Step-by-Step Deployment Procedure

### 1. Install Docker and Docker Compose.

```
# Debian/Ubuntu
apt-get install -y docker.io docker-compose-v2
systemctl enable --now docker
```

### 2. Clone the repository and build the image.

```
git clone https://gitlab.example.com/porr-ag/infra-webshop.git
cd infra-webshop
make docker-build
# or pull from registry:
docker login registry.gitlab.example.com
docker pull
    registry.gitlab.example.com/porr-ag/infra-webshop:latest
```

### 3. Create the environment file.

```
cp deploy/docker-host/.env.example deploy/docker-host/.env
# Edit .env with a text editor; set at minimum:
#   DATABASE_URL, SESSION_KEY, ADMIN_EMAIL, ADMIN_PASSWORD,
#   BASE_URL
```

### 4. Start PostgreSQL and the application.

```
cd deploy/docker-host
docker compose up -d
```

### 5. Run database migrations.

```
docker exec infra-webshop /migrate
```

## 6. Obtain a TLS certificate.

*Option A — Let's Encrypt (public domain):*

```
apt-get install -y certbot python3-certbot-nginx
certbot --nginx -d shop.example.com
```

*Option B — Corporate CA (internal PKI):* Place the signed certificate and key at the paths referenced in [deploy/docker-host/nginx.conf](#) (see below) and restart Nginx.

## 7. Configure Nginx as TLS-terminating reverse proxy.

Listing 6.1 shows a minimal but production-ready Nginx configuration:

```
server {
    listen 80;
    server_name shop.example.com;
    return 301 https://$host$request_uri;
}

server {
    listen 443 ssl http2;
    server_name shop.example.com;

    ssl_certificate
        /etc/letsencrypt/live/shop.example.com/fullchain.pem;
    ssl_certificate_key
        /etc/letsencrypt/live/shop.example.com/privkey.pem;
    ssl_protocols      TLSv1.2 TLSv1.3;
    ssl_ciphers         HIGH:!aNULL:!MD5;

    client_max_body_size 10m;

    location / {
        proxy_pass      http://127.0.0.1:8080;
        proxy_set_header Host          $host;
        proxy_set_header X-Real-IP     $remote_addr;
        proxy_set_header X-Forwarded-For
            $proxy_add_x_forwarded_for;
        proxy_set_header X-Forwarded-Proto $scheme;
        proxy_read_timeout 60s;
    }
}
```

Listing 6.1: `deploy/docker-host/nginx.conf`

## 8. Reload Nginx.

```
nginx -t && systemctl reload nginx
```

## 9. Verify the deployment.

Navigate to <https://shop.example.com> and log in with **ADMIN\_EMAIL** / **ADMIN\_PASSWORD**. Change the password immediately under **Settings** → **My Profile**.

#### 10. (Optional) Enable automatic certificate renewal.

```
# Certbot installs a systemd timer automatically
systemctl status certbot.timer
```

## Backup & Restore

```
# Backup PostgreSQL
docker exec porr-postgres pg_dump -U webshop webshop \
  | gzip > backup-$(date +%Y%m%d).sql.gz

# Restore
zcat backup-20260101.sql.gz \
  | docker exec -i porr-postgres psql -U webshop webshop
```

## 6.4 Deployment on Kubernetes

This section covers deploying the webshop to an existing Kubernetes cluster (e.g. the LKE cluster provisioned by k8s-platform).

### Architecture




Figure 6.2: Kubernetes deployment architecture

### Prerequisites

- **kubectl** and **helm** (v3.12+) configured for the target cluster
- Nginx Ingress Controller installed (**ingress-nginx**)
- **cert-manager** installed with a **ClusterIssuer** for Let's Encrypt or a corporate CA
- The image **porrag/infra-webshop:tag** is publicly available on Docker Hub — no **imagePullSecret** required

## Deployment with Helm

The repository ships a Helm chart at [deploy/helm/infra-webshop/](#). It covers the full deployment in a single command and handles database migrations automatically as a pre-install/pre-upgrade hook.

```
helm upgrade --install infra-webshop \
  deploy/helm/infra-webshop \
  --namespace infra-webshop \
  --create-namespace \
  --set
    env.databaseUrl="postgres://user:pass@pg.example.com:5432/webshop?sslmode=require" \
  --set env.sessionSecret="<32-byte-hex>" \
  --set env.adminPassword="<initial-password>" \
  --set ingress.hosts[0].host=shop.example.com \
  --set ingress.tls[0].hosts[0]=shop.example.com \
  --set ingress.tls[0].secretName=infra-webshop-tls
```

Listing 6.2: Helm install with required values

### Note

Sensitive values should be passed via a separate [secrets.values.yaml](#) file (not committed to version control) rather than as `-set` arguments, so they do not appear in shell history.

```
env:
  databaseUrl:
    "postgres://user:pass@pg.example.com:5432/webshop?sslmode=require"
  sessionSecret:
    "aabbccddeeff00112233445566778899aabbccddeeff00112233445566778899"
  adminPassword: "change-me-immediately"
  entra:
    tenantId: "xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx"
    clientId: "xxxxxxxx-xxxx-xxxx-xxxx-xxxxxxxxxxxx"
    clientSecret: "your-client-secret"
    redirectUrl: "https://shop.example.com/auth/callback"
```

Listing 6.3: Example secrets.values.yaml (do not commit)

```
helm upgrade --install infra-webshop deploy/helm/infra-webshop \
  --namespace infra-webshop --create-namespace \
  -f values.production.yaml \
  -f secrets.values.yaml
```

Table 6.2 lists the most important configurable values.

Table 6.2: Key Helm values

| Key                     | Description                    | Default              |
|-------------------------|--------------------------------|----------------------|
| image.repository        | Docker Hub image name          | porrag/infra-webshop |
| image.tag               | Image tag (empty = appVersion) | ""                   |
| replicaCount            | Deployment replicas            | 2                    |
| env.databaseUrl         | PostgreSQL DSN (required)      | —                    |
| env.sessionSecret       | Session encryption key (req.)  | —                    |
| env.adminPassword       | Initial admin password (req.)  | —                    |
| ingress.hosts           | Ingress host list              | shop.example.com     |
| ingress.tls             | TLS configuration              | (see values.yaml)    |
| autoscaling.enabled     | Enable HPA                     | false                |
| autoscaling.maxReplicas | HPA max replica count          | 6                    |
| migrate.enabled         | Run DB migrations as hook      | true                 |

## Step-by-Step Manual Deployment (without Helm)

### 1. Create the namespace.

```
kubectl create namespace infra-webshop
```

### 2. Create the image pull secret (if using a private registry).

```
kubectl create secret docker-registry gitlab-registry \
  --namespace infra-webshop \
  --docker-server=registry.gitlab.example.com \
  --docker-username=<deploy-token-name> \
  --docker-password=<deploy-token-value>
```

### 3. Create the application Secret.

```
apiVersion: v1
kind: Secret
metadata:
  name: infra-webshop-env
  namespace: infra-webshop
type: Opaque
stringData:
  DATABASE_URL:
    "postgres://webshop:password@pg.example.com:5432/webshop?sslmode=require"
  SESSION_KEY:
    "000102030405060708090a0b0c0d0e0f101112131415161718191a1b1c1d1e1f"
  ADMIN_EMAIL: "admin@example.com"
  ADMIN_PASSWORD: "changeme"
  BASE_URL: "https://shop.example.com"
```

Listing 6.4: Kubernetes Secret for the webshop

### 4. Run database migrations as a Kubernetes Job.

```
apiVersion: batch/v1
kind: Job
metadata:
  name: infra-webshop-migrate
  namespace: infra-webshop
spec:
  template:
    spec:
      restartPolicy: Never
      imagePullSecrets:
        - name: gitlab-registry
      containers:
        - name: migrate
          image: registry.gitlab.example.com/porr-ag/infra-webshop:latest
          command: ["/migrate"]
          envFrom:
            - secretRef:
                name: infra-webshop-env
```

Listing 6.5: Migration Job

```
kubectl apply -f migrate-job.yaml
kubectl wait --for=condition=complete
  job/infra-webshop-migrate \
  -n infra-webshop --timeout=120s
```

## 5. Deploy the application.

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: infra-webshop
  namespace: infra-webshop
spec:
  replicas: 2
  selector:
    matchLabels:
      app: infra-webshop
  template:
    metadata:
      labels:
        app: infra-webshop
    spec:
      imagePullSecrets:
        - name: gitlab-registry
      containers:
        - name: app
          image: registry.gitlab.example.com/porr-ag/infra-webshop:latest
          ports:
            - containerPort: 8080
          envFrom:
            - secretRef:
                name: infra-webshop-env
```

```

        livenessProbe:
          httpGet:
            path: /health
            port: 8080
          initialDelaySeconds: 10
          periodSeconds: 30
        readinessProbe:
          httpGet:
            path: /health
            port: 8080
          initialDelaySeconds: 5
          periodSeconds: 10
        resources:
          requests:
            cpu: "100m"
            memory: "128Mi"
          limits:
            cpu: "500m"
            memory: "512Mi"
---
apiVersion: v1
kind: Service
metadata:
  name: infra-webshop
  namespace: infra-webshop
spec:
  selector:
    app: infra-webshop
  ports:
    - port: 80
      targetPort: 8080

```

Listing 6.6: Deployment and Service

## 6. Configure the Ingress with TLS.

```

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: infra-webshop
  namespace: infra-webshop
  annotations:
    nginx.ingress.kubernetes.io/proxy-body-size: "10m"
    cert-manager.io/cluster-issuer: "letsencrypt-prod"
spec:
  ingressClassName: nginx
  tls:
    - hosts:
        - shop.example.com
      secretName: infra-webshop-tls
  rules:
    - host: shop.example.com
      http:
        paths:
          - path: /
            pathType: Prefix

```

```
backend:
  service:
    name: infra-webshop
    port:
      number: 80
```

Listing 6.7: Ingress with cert-manager TLS

## 7. (Optional) Configure Horizontal Pod Autoscaler.

```
apiVersion: autoscaling/v2
kind: HorizontalPodAutoscaler
metadata:
  name: infra-webshop
  namespace: infra-webshop
spec:
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: infra-webshop
  minReplicas: 2
  maxReplicas: 6
  metrics:
    - type: Resource
      resource:
        name: cpu
        target:
          type: Utilization
          averageUtilization: 70
```

Listing 6.8: HorizontalPodAutoscaler

## 8. Apply all manifests and verify.

```
kubectl apply -f secret.yaml
kubectl apply -f deployment.yaml
kubectl apply -f ingress.yaml
kubectl apply -f hpa.yaml          # optional

kubectl rollout status deployment/infra-webshop -n
  infra-webshop
kubectl get ingress -n infra-webshop
```

### Note

The polling goroutines that track GitLab pipeline status run inside each pod. With multiple replicas, advisory locks in PostgreSQL (`pg_try_advisory_lock`) ensure that only one pod polls each pipeline at a time, preventing duplicate status updates.

## 6.5 Database Migrations

Migrations are versioned SQL files in `internal/migrations/`. Run migrations before starting the server:

```
# Run all pending migrations
make migrate

# Or directly
go run ./cmd/migrate
```

Migrations are idempotent (IF NOT EXISTS, IF NOT EXISTS guards). Re-running is safe.

## 6.6 First Login

1. Navigate to the webshop URL.
2. Click **Local Login** (not **Sign in with Microsoft**).
3. Enter the email and password from **ADMIN\_EMAIL** and **ADMIN\_PASSWORD**.
4. Navigate to **Settings** → **My Profile** and change the default password immediately.

# Chapter 7

## System Configuration

### 7.1 GitLab Sources

Under **Administration** → **GitLab Sources**:

1. Click **Add New Source**.
2. Enter a **Name** (label), the instance **URL**, and a **Personal Access Token** with `read_api` scope.
3. Save. The source is now available when creating environments and when browsing repositories for parameter import.

### 7.2 Deployment Environments

Under **Administration** → **Deployment Environments**:

1. Click **Add New Environment**.
2. Enter a **Name** (e.g. *Linode Frankfurt Production*), select the **GitLab Source**, and enter the **Webhook URL** and **Webhook Token**.
3. Save. The environment becomes available when assigning products.

#### Note

The Webhook URL points to a specific GitLab project pipeline trigger endpoint, e.g.:

`https://gitlab.example.com/api/v4/projects/42/ref/main/trigger/pipeline`

### 7.3 SMTP Configuration

Under **Administration** → **Email** (</admin/smtp>):

Enter host, port, sender address, username, password, and TLS mode. Click **Save**. The notification service reloads credentials immediately without a server restart.

## 7.4 AI Translation

Under **Administration** → **AI Translation** (</admin/ai-config>):

1. Select a **Provider**: Claude, OpenAI, Azure OpenAI, Ollama, or LocalAI.
2. Enter the **API endpoint** and **API key**. For Ollama/LocalAI, the endpoint is a local URL and no key is required.
3. Enter the **model name** (e.g. `claude-opus-4-7`, `gpt-4o`, `llama3.1`).
4. Click **Save**. The translator instance is replaced immediately — the next translation request uses the new provider.

## 7.5 Branding

Under **Administration** → **Shop Design** (</admin/branding>):

- **Shop Name** and **Subtitle** — shown in the header.
- **Primary Color** — header, footer, navigation (default: `#131921`).
- **Accent Color** — buttons and call-to-action elements (default: `#febd69`).
- **Logo** — PNG or SVG, displayed in the header; replaces the shop name text if present.
- **Imprint** — HTML-formatted legal imprint text, shown at </impressum>.

All branding values are stored in a singleton database row and cached in-memory with a mutex for read performance.

## 7.6 Currencies

Under **Administration** → **Currencies**:

1. The **base currency** (default EUR) is the currency in which all product prices are stored.
2. Optionally configure an exchange rate API key and endpoint.
3. Click **Update Rates** to fetch live rates.
4. Individual rates can be overridden manually for fixed-rate environments.

Prices are displayed in the user's locale currency by looking up the browser locale and applying the stored rate.

# Chapter 8

## User Management

### 8.1 Roles

Table 8.1: User roles and capabilities

| Role            | Description                            | Auth method | DB role    |
|-----------------|----------------------------------------|-------------|------------|
| DU Admin        | Full access; direct ordering; approval | SSO         | admin      |
| Webshop Admin   | Catalog & config; no ordering          | Local       | shop_admin |
| Project Manager | Orders own infra; approval required    | SSO         | user       |

### 8.2 Creating a Local Account

Under **Administration** → **Users**:

1. Click **New User**.
2. Enter **Name**, **Email**, **Password**, and select a **Role**.
3. Click **Save**.

SSO users are created automatically on first login.

### 8.3 Editing Users

Click **Edit** next to a user to change name, email, or role. The password can only be changed by the user themselves under **Settings** → **My Profile**.

## 8.4 Deactivating Users

Click **Deactivate** on the user edit page. Deactivated users cannot log in. Their orders and projects remain visible in the system.

# Chapter 9

## Audit Log

The audit log under `/audit` is an append-only, tamper-evident record of all system actions.

### 9.1 Logged Events

Table 9.1: Audited actions

| Action                | Trigger                         |
|-----------------------|---------------------------------|
| order_placed          | User submits an order           |
| order_approved        | DU Admin approves               |
| order_rejected        | DU Admin rejects (with comment) |
| order_deployed        | Pipeline completes successfully |
| order_failed          | Pipeline fails                  |
| decommission_started  | User triggers decommission      |
| decommission_complete | Destroy pipeline succeeds       |
| user_login            | Successful authentication       |
| user_created          | Admin creates a user            |
| config_changed        | System configuration updated    |

### 9.2 Filtering

Use the filter form to narrow results by:

- **Date range** (from / to)
- **User**
- **Action type**

## 9.3 Export

Click **Export CSV** or **Export PDF** to download the currently filtered audit log for compliance reporting.

# Part IV

## Developer Guide

# Chapter 10

## Development Environment

### 10.1 Prerequisites

- Go 1.22+
- Node.js 20+ (for Tailwind CSS compilation)
- Docker & Docker Compose (for local PostgreSQL)
- `pdflatex` or `lualatex` (for handbook compilation)

### 10.2 Quick Start

```
# Clone and enter the repo
git clone git@gitlab.example.com:porr-ag/infra-webshop.git
cd infra-webshop

# Copy and populate environment
cp .env.example .env
# Edit .env with your local DB credentials

# Start PostgreSQL and the server in one command
make dev

# In another terminal: CSS hot-reload
make css-watch
```

`make dev` starts Docker Compose (PostgreSQL), runs CSS build, then starts the Go server with live environment variables from `.env`.

## 10.3 Makefile Targets

Table 10.1: Makefile targets

| Target                         | Action                             |
|--------------------------------|------------------------------------|
| <code>make build</code>        | Build CSS + Go binary              |
| <code>make run</code>          | Start server (requires running DB) |
| <code>make migrate</code>      | Run database migrations            |
| <code>make css</code>          | Build Tailwind CSS once            |
| <code>make css-watch</code>    | Watch and rebuild CSS on change    |
| <code>make test</code>         | Run all Go tests                   |
| <code>make vet</code>          | Run <code>go vet</code>            |
| <code>make docker-build</code> | Build Docker image                 |
| <code>make dev</code>          | Start DB + server (development)    |
| <code>make dev-down</code>     | Stop DB + server                   |
| <code>make clean</code>        | Remove build artefacts             |
| <code>make docs</code>         | Compile the LaTeX handbook to PDF  |

# Chapter 11

## Codebase Structure

### 11.1 Directory Layout

```
infra-webshop/
+-- cmd/
|   +-- server/          # main.go -- server entry point
|   +-- migrate/         # main.go -- migration runner
+-- docs/                # Documentation (this handbook)
|   +-- handbook.tex
|   +-- handbook.pdf     # compiled output
+-- internal/
|   +-- aitranslation/   # AI translator interface + backends
|   +-- audit/           # AuditService implementation
|   +-- auth/            # OIDC, session store, password hashing
|   +-- config/          # Config struct loaded from env
|   +-- exchange/        # Currency exchange rate service
|   +-- handler/         # HTTP handlers and route registration
|   +-- i18n/            # 25-language key-value translation map
|   +-- migrations/      # SQL migration files (001_*.sql, ...)
|   +-- model/           # Shared domain model structs
|   +-- notification/    # SMTP notification service
|   +-- repository/      # Repository interfaces
|   |   +-- postgres/    # pgx/v5 implementations
|   +-- service/         # Service interfaces
|   +-- impl/            # Service implementations
+-- ui/
|   +-- templates/
|   |   +-- layout.html  # Shared layout
|   |   +-- pages/       # One file per page
|   |   +-- partials/    # HTMX fragment templates
|   +-- static/
|   |   +-- css/style.css # Compiled Tailwind output
+-- ui.go                # embed.FS declarations
```

### 11.2 Layered Architecture

The code follows a strict layered architecture with downward-only dependencies:

Handler → Service → Repository → PostgreSQL

**Handler layer** ([internal/handler/](#)) HTTP handlers only. Read request, call service, render template. No business logic. No SQL.

**Service layer** ([internal/service/](#)) Business rules and orchestration. Calls repositories. Triggers notifications, audit writes, webhook dispatches. Defined as interfaces in [service.go](#), implemented in [impl/](#).

**Repository layer** ([internal/repository/](#)) Data access. Defined as interfaces (one per aggregate root). Implemented in [postgres/](#) using pgx. No business logic.

**Model layer** ([internal/model/](#)) Plain Go structs shared across layers. No methods beyond simple helpers.

## 11.3 Template System

Each page is a separate `*template.Template` instance containing the shared [layout.html](#) plus its own page file. This avoids `{{define "content"}}` collisions between pages.

Page templates define two blocks:

```
{{define "title"}}Page Title{{end}}
{{define "content"}}
  <!-- page HTML using Tailwind / DaisyUI -->
{{end}}
```

Data is passed as `map[string]any` to `h.render()`, which wraps it in a `PageData` struct providing:

- .Lang** User's active language code (e.g. "de").
- .Brand** Branding values (colors, shop name, etc.).
- .Session** Current user session (name, role, email).
- .Flash** One-time flash message (success/error).
- .Data** The `map[string]any` passed by the handler.

Access data in templates:

```
{{t "admin.title" .Lang}}      {{/* translated string */}}
{{.Brand.PrimaryColor}}        {{/* branding color */}}
{{.Session.Name}}              {{/* logged-in user's name */}}
{{with .Data}}
  {{.ProductCount}}            {{/* handler-provided data */}}
{{end}}
```

## 11.4 Adding Translations

All UI strings live in `internal/i18n/i18n.go` as a `map[string]map[string]string`. To add a new key:

```
// In translations map:
"my.new.key": {
    "en": "English text",
    "de": "Deutscher Text",
    "fr": "Texte francais",
    // ... all 25 language codes
},
```

Use in templates with `{{t "my.new.key" .Lang}}`.

## 11.5 Database Migrations

Migration files follow the naming pattern `NNN_description.sql` and are executed in alphabetical order. Write migrations defensively:

```
ALTER TABLE products
  ADD COLUMN IF NOT EXISTS image_mime TEXT NOT NULL DEFAULT
    'image/jpeg';

CREATE TABLE IF NOT EXISTS app_config (
  id INT PRIMARY KEY DEFAULT 1 CHECK (id = 1),
  smtp_host TEXT NOT NULL DEFAULT ''
  -- ...
);
```

### Important

Never modify a migration that has already been applied to production. Always create a new migration file for schema changes.

# Chapter 12

## Adding a New Feature

This chapter describes the canonical pattern for adding a feature end-to-end.

### 12.1 Example: Adding a New Admin Page

**Goal:** Add an admin page at `/admin/widgets` that lists and creates widgets stored in a new `widgets` database table.

#### 12.1.1 Step 1: Create the Migration

Create `internal/migrations/010_widgets.sql`:

```
CREATE TABLE IF NOT EXISTS widgets (  
  id    BIGSERIAL PRIMARY KEY,  
  name  TEXT NOT NULL,  
  active BOOLEAN NOT NULL DEFAULT TRUE,  
  created_at TIMESTAMPTZ NOT NULL DEFAULT now()  
);
```

#### 12.1.2 Step 2: Add the Model

In `internal/model/model.go`:

```
type Widget struct {  
  ID          int64  
  Name        string  
  Active      bool  
  CreatedAt   time.Time  
}
```

#### 12.1.3 Step 3: Define the Repository Interface

In `internal/repository/repository.go`:

```
type WidgetRepository interface {
    FindAll(ctx context.Context) ([]model.Widget, error)
    Create(ctx context.Context, name string) error
    Delete(ctx context.Context, id int64) error
}
```

#### 12.1.4 Step 4: Implement the Repository

Create `internal/repository/postgres/widget.go`:

```
type widgetRepo struct{ pool *pgxpool.Pool }

func NewWidgetRepository(pool *pgxpool.Pool)
    repository.WidgetRepository {
    return &widgetRepo{pool: pool}
}

func (r *widgetRepo) FindAll(ctx context.Context) ([]model.Widget,
    error) {
    rows, err := r.pool.Query(ctx,
        'SELECT id, name, active, created_at FROM widgets ORDER BY
        name')
    // ... scan rows into []model.Widget
}
```

#### 12.1.5 Step 5: Add the Handler

Add to the Handler struct in `routes.go`:

```
widgets repository.WidgetRepository
```

Add to Deps struct and wire in `New()`. Add handler methods to a new file `handler/widgets.go`:

```
func (h *Handler) adminWidgets(w http.ResponseWriter, r
    *http.Request) {
    items, _ := h.widgets.FindAll(r.Context())
    h.render(w, r, "admin-widgets.html", map[string]any{
        "Widgets": items,
    })
}
```

Register routes in `RegisterRoutes`:

```
mux.Handle("GET /admin/widgets", admin(h.adminWidgets))
mux.Handle("POST /admin/widgets", admin(h.adminWidgetCreate))
```

#### 12.1.6 Step 6: Create the Template

Create `ui/templates/pages/admin-widgets.html`:

```
{{define "title"}}{{t "admin.widgets" .Lang}}{{end}}
{{define "content"}}
<div class="max-w-4xl mx-auto">
  <h2 class="text-2xl font-bold mb-6">{{t "admin.widgets"
    .Lang}}</h2>
  {{with .Data}}
  <table class="table table-zebra w-full">
    {{range .Widgets}}
    <tr><td>{{.Name}}</td></tr>
    {{end}}
  </table>
  {{end}}
</div>
{{end}}
```

### 12.1.7 Step 7: Add i18n Keys

In [internal/i18n/i18n.go](#), add the key with all 25 language translations.

### 12.1.8 Step 8: Wire Up in main.go

```
widgetRepo := pgRepo.NewWidgetRepository(pool)
// ...
h := handler.New(handler.Deps{
  // ...
  Widgets: widgetRepo,
})
```

### 12.1.9 Step 9: Test

Add test stubs in [internal/repository/](#) (or use the real DB in integration tests).  
Run:

```
make test
make vet
```

## Part V

# Project Manager & Product Lifecycle Guide

# Chapter 13

## Ordering Infrastructure

### 13.1 The Order Lifecycle




Figure 13.1: Order lifecycle state machine

### 13.2 Placing an Order

1. Navigate to **Catalog** and select a product.
2. Click **Configure** → select the deployment environment.
3. Select or create a **Project**.
4. Fill in the **Parameters** (required fields are marked).
5. Select or confirm the **Cost Center**.
6. Click **Deploy Now**.

As a **Project Manager** the order status changes to *Pending Approval*. A DU Admin receives an email notification.

As a **DU Admin** the GitLab webhook fires immediately.

## 13.3 Tracking Status

After placing an order, navigate to **My Orders** → **{order}**. The status updates live via HTMX polling every 30 seconds.

## 13.4 Using an Existing Deployment as a Template

1. Open **Infrastructure**.
2. Find the resource and click **Use as Template**.
3. The order form opens pre-filled with the original parameters.
4. Modify as needed and submit.

## 13.5 Decommissioning

1. Open **Infrastructure**.
2. Find the resource and click **Decommission**.
3. Confirm the action.
4. The destroy webhook fires; status changes to *Decommissioning*.
5. On pipeline completion, status becomes *Decommissioned*.

# Chapter 14

## Introducing a New Product

This chapter describes the complete process for making a new infrastructure product available in the webshop. Follow the steps sequentially.

### 14.1 Overview

1. Plan the product
2. Prepare the OpenTofu module
3. Configure the GitLab webhook
4. Create the product in the Admin UI
5. Define parameters (import or manual)
6. Configure deployment environments and pricing
7. Set up AI translations
8. Test the full order flow
9. Publish

### 14.2 Step 1: Plan the Product

Before touching any UI, answer these questions:

- **What does the product deploy?** (e.g. *Managed PostgreSQL VM on Linode*)
- **Which OpenTofu module delivers it?** (vm-platform, k8s-platform, or a new module)
- **In which environments is it available?** (e.g. Linode Frankfurt, Linode Amsterdam)
- **What parameters does the user need to provide?** (e.g. `vm_label`, `instance_type`, `db_name`)

- **What is the base price in EUR per environment?**
- **Who assigns the cost center?** (project / select / overhead)

## 14.3 Step 2: Prepare the OpenTofu Module

If the product uses an existing module (vm-platform or k8s-platform), skip to Step 3. For a new product type:

1. Create a new GitLab repository in the `porr-ag` group.
2. Add the standard files: `provider.tf`, `backend.tf`, `variables.tf`, `main.tf`, `outputs.tf`, `.gitlab-ci.yml`.
3. Declare all user-provided parameters as `variable` blocks in `variables.tf` with `description`, `type`, and `default` where applicable:

```
variable "db_name" {
  description = "Name of the PostgreSQL database to create."
  type        = string
}

variable "db_size_gb" {
  description = "Database volume size in GB."
  type        = number
  default     = 20
}
```

4. Test the module locally:

```
tofu init -backend-config="address=..." \
          -backend-config="username=..." \
          -backend-config="password=..."
tofu plan -var="linode_token=$TOKEN" \
          -var="db_name=test" \
          -var="db_size_gb=20"
tofu apply
tofu destroy
```

5. Commit and push to the `main` branch.

## 14.4 Step 3: Configure the GitLab Webhook

In the GitLab repository:

1. Go to **Settings** → **CI/CD** → **Pipeline trigger tokens**.
2. Click **Add new token** and give it a descriptive name (e.g. *webshop-prod*).
3. Note the **trigger URL** and **token**.

The trigger URL has the form:

```
https://gitlab.example.com/api/v4/projects/<PROJECT_ID>/ref/main/trigger/pipeline
```

## 14.5 Step 4: Create a Category (if needed)

Under **Administration** → **Categories**:

1. Click **New Category**.
2. Enter a **Name**.
3. Optionally assign a **Category Parameter Set** (see Step 5).
4. Save.

## 14.6 Step 5: Create the Product

Under **Administration** → **Products** → **New**:

### 14.6.1 Basic Information

1. Select the **Category**.
2. Enter the product **Name** in the primary language.
3. Enter a detailed **Description** in the primary language.
4. Upload a product **Image** (PNG or JPEG, max 10 MB).
5. Click **Save & Continue**.

### 14.6.2 Parameters

Choose one of two import methods:

**Option A — Import from `variables.tf`** (recommended):

1. Click **Browse Repository**.
2. Select the **GitLab Source** and the repository.
3. Select the branch (`main`) and navigate to `variables.tf`.
4. Click **Import Parameters**.
5. The webshop parses the HCL and creates parameter records for each variable.  
Review and adjust:
  - **Required** — must the user fill this in?
  - **Sensitive** — should the value be masked in the UI and audit log?
  - **Default value** — pre-fill the order form.

- **Type** — text, number, boolean, or dropdown.

**Option B — Manual entry:**

Click **Add Parameter** and fill in each field manually.

**Note**

Parameters whose names match OpenTofu variable names (e.g. `instance_type`) are passed to the webhook as `TF_VAR_instance_type`. The name in the webshop must match the `variable` block name in `variables.tf` exactly.

### 14.6.3 Deployment Environments

1. Click **Add Environment**.
2. Select an environment from the dropdown (must be pre-configured under Administration → Environments).
3. Enter the **Price** in the base currency.
4. For the webhook, either:
  - Use the environment's default webhook URL, or
  - Enter a product-specific webhook URL and token.
5. Configure the **Cost Center Mode**:
 

**Project** Cost center inherited from the project.

**Select** Orderer selects from the cost center list.

**Overhead** Always assigns a fixed cost center.
6. Repeat for each environment.

### 14.6.4 Multiple Webhooks (Pipeline Arrays)

A single product-environment combination can have multiple webhooks executed in order (e.g. first provision the VM, then configure DNS, then register in Vault). Under the webhook section:

1. Click **Add Webhook**.
2. Enter URL, token, and **Execution Order** (1, 2, 3 ...).
3. Webhooks are triggered sequentially in ascending order.

## 14.7 Step 6: Configure AI Translation

1. Ensure an AI provider is configured (Chapter 7).
2. Open the product in **Administration** → **Products**.

3. Click **Generate AI Translation**.
4. Review the generated translations for all 25 languages.
5. Edit any that require adjustment.
6. Save.

#### Tip

Generate translations early in the process. The AI translates both the product name and description. Individual translations can be re-generated for specific languages by editing the translation tab directly.

## 14.8 Step 7: Global and Category Parameter Sets

**Global parameters** (under **Administration** → **Global Parameters**) apply to every order regardless of product. Use them for organization-wide fields such as a cost center code or project tag that must always be captured.

**Category parameters** (under **Administration** → **Categories** → **{category}** → **Parameters**) apply to all products within a category. Use them to avoid repeating parameters shared by related products (e.g. all VM products share a **region** parameter).

**Parameter hierarchy** (applied in this order on the order form):

1. Global parameters
2. Category parameters
3. Product parameters
4. Environment-specific parameters (override per-environment)

## 14.9 Step 8: Test the Full Order Flow

1. Log in as a **DU Admin** (DU Admins trigger webhooks directly).
2. Navigate to **Catalog** and find the new product.
3. Click **Configure**, fill in test parameters, and click **Deploy Now**.
4. Monitor the order under **My Orders** — the status should progress from *Provisioning* to *Completed*.
5. Verify the infrastructure was created in the Linode Cloud Manager.
6. Test the email notification (check the admin email inbox).
7. Test the **Decommission** flow:
  - Open **Infrastructure**.
  - Click **Decommission** on the test resource.

- Verify the destroy pipeline runs and the resource disappears from the Linode Cloud Manager.

**Important**

Always test decommission in a non-production environment first. A destroy pipeline permanently deletes cloud resources and their data.

## 14.10 Step 9: Test with a Project Manager Account

1. Log in (or create a test account) as a **Project Manager**.
2. Place an order for the new product.
3. Verify the order shows *Pending Approval*.
4. Switch to the **DU Admin** account.
5. Approve the order under **Approvals**.
6. Verify the webhook fires and provisioning completes.
7. Verify the Project Manager receives the confirmation email.

## 14.11 Step 10: Product is Live

Once testing passes:

- The product is immediately visible in the catalog.
- No deployment or restart of the webshop is required.
- Announce the new product to users via the notification service or the shop's subtitle/branding.

# Chapter 15

## Project Management

### 15.1 Projects

A *project* is a container for related infrastructure orders. Every order must be associated with a project. Projects have:

- A **Name** and optional **Description**.
- An optional **Cost Center** (used when cost center mode is *Project*).
- An owner (the user who created it).

**Project Managers** see only their own projects. **DU Admins** and **Webshop Admins** see all projects.

### 15.2 Infrastructure Overview

The **Infrastructure** page shows all deployed resources grouped by project and environment. Each resource shows:

- Product name and environment.
- Deployment date and current status.
- Key parameters (e.g. VM label, cluster name).
- Buttons: **Use as Template**, **Decommission**.

### 15.3 Cost Center Assignment

Cost centers are maintained under **Administration** → **Cost Centers**. The three assignment modes give fine-grained control over who can allocate costs:

**Project** No user input; cost center comes from the project. Use for products where cost is always allocated to the project budget.

**Select** Orderer picks from the cost center list. Use when individual orders can span cost centers.

**Overhead** Always assigns a fixed cost center regardless of order. Use for shared services (monitoring, DNS).

# Appendix A

## Supported Languages

The following 25 languages are supported in the UI and for AI product translation:

| Code | Language   | Native name             |
|------|------------|-------------------------|
| de   | German     | Deutsch                 |
| en   | English    | English                 |
| fr   | French     | Français                |
| it   | Italian    | Italiano                |
| es   | Spanish    | Español                 |
| pt   | Portuguese | Português               |
| nl   | Dutch      | Nederlands              |
| pl   | Polish     | Polski                  |
| cs   | Czech      | Čeština                 |
| sk   | Slovak     | Slovenčina              |
| hu   | Hungarian  | Magyar                  |
| ro   | Romanian   | Română                  |
| bg   | Bulgarian  | Balgarski (Cyrillic)    |
| hr   | Croatian   | Hrvatski                |
| sl   | Slovenian  | Slovenščina             |
| et   | Estonian   | Eesti                   |
| lv   | Latvian    | Latviešu                |
| lt   | Lithuanian | Lietuvių                |
| fi   | Finnish    | Suomi                   |
| sv   | Swedish    | Svenska                 |
| da   | Danish     | Dansk                   |
| el   | Greek      | Ellinika (Greek script) |
| mt   | Maltese    | Malti                   |
| ru   | Russian    | Russkij (Cyrillic)      |
| uk   | Ukrainian  | Ukrayinska (Cyrillic)   |

Table A.1: Supported language codes

# Appendix B

## Environment Variable Reference

| Variable           | Description                                                    | Required |
|--------------------|----------------------------------------------------------------|----------|
| DATABASE_URL       | PostgreSQL connection string (DSN)                             | Yes      |
| SESSION_KEY        | 32-byte hex key for session encryption                         | Yes      |
| ADMIN_EMAIL        | Email of the initial Webshop Admin account                     | Yes      |
| ADMIN_PASSWORD     | Password of the initial admin account                          | Yes      |
| BASE_URL           | Public base URL (e.g. <code>https://shop.example.com</code> )  | Yes      |
| PORT               | HTTP listen port (default: 8080)                               | No       |
| OIDC_CLIENT_ID     | Entra ID application (client) ID                               | No       |
| OIDC_CLIENT_SECRET | Entra ID client secret                                         | No       |
| OIDC_ISSUER        | OIDC issuer URL<br>( <code>https://login.microsoft...</code> ) | No       |
| SMTP_HOST          | SMTP server hostname                                           | No       |
| SMTP_PORT          | SMTP port (default: 587)                                       | No       |
| SMTP_FROM          | Sender email address                                           | No       |
| SMTP_USERNAME      | SMTP authentication username                                   | No       |
| SMTP_PASSWORD      | SMTP authentication password                                   | No       |
| SMTP_TLS           | <code>true</code> to use TLS/SSL                               | No       |
| AI_PROVIDER        | AI translation provider at startup                             | No       |
| AI_ENDPOINT        | AI API endpoint URL                                            | No       |
| AI_API_KEY         | AI API key                                                     | No       |
| AI_MODEL           | AI model name                                                  | No       |
| EXCHANGE_RATE_KEY  | Exchange rate API key                                          | No       |

Table B.1: Environment variable reference

# Appendix C

## OpenTofu Module Variable Reference

### C.1 network-core Variables

| Variable               | Type      | Description                  | Default                  |
|------------------------|-----------|------------------------------|--------------------------|
| linode_token           | string    | Linode API token (sensitive) | —                        |
| region                 | string    | Linode region                | eu-central               |
| vpc_label              | string    | VPC label                    | porr-vpc-prod            |
| vpc_description        | string    | VPC description              | (see defaults)           |
| subnets                | map(obj)  | Subnet map (label→cidr)      | app + mgmt               |
| firewall_inbound_rules | list(obj) | Inbound firewall rules       | SSH + ICMP               |
| dns_domain             | string    | DNS zone domain (empty=none) | ""                       |
| tags                   | list      | Resource tags                | ["infra","network-core"] |

Table C.1: network-core variables

### C.2 vm-platform Variables

| Variable                 | Type      | Description                  | Default         |
|--------------------------|-----------|------------------------------|-----------------|
| linode_token             | string    | Linode API token (sensitive) | —               |
| vm_label                 | string    | Unique VM label (state name) | —               |
| region                   | string    | Linode region                | eu-central      |
| vpc_id                   | number    | VPC ID from network-core     | —               |
| subnet_id                | number    | Subnet ID from network-core  | —               |
| firewall_id              | number    | Base firewall ID             | —               |
| instance_type            | string    | Linode plan                  | g6-standard-2   |
| image                    | string    | Boot image                   | linode/debian12 |
| root_pass                | string    | Root password (sensitive)    | —               |
| ssh_keys                 | list      | Authorized SSH keys          | []              |
| volume_size_gb           | number    | Block storage size (0=none)  | 0               |
| additional_inbound_rules | list(obj) | Extra firewall rules         | []              |

| Variable     | Type   | Description             | Default        |
|--------------|--------|-------------------------|----------------|
| dns_zone_id  | number | DNS zone ID (null=none) | null           |
| dns_hostname | string | A record hostname       | ""             |
| tags         | list   | Resource tags           | ["infra","vm"] |

Table C.2: vm-platform variables

## C.3 k8s-platform Variables

| Variable                 | Type      | Description                  | Default          |
|--------------------------|-----------|------------------------------|------------------|
| linode_token             | string    | Linode API token (sensitive) | —                |
| cluster_label            | string    | Unique cluster label         | —                |
| region                   | string    | Linode region                | eu-central       |
| kubernetes_version       | string    | K8s version                  | 1.31             |
| vpc_id                   | number    | VPC ID                       | —                |
| subnet_id                | number    | Subnet ID                    | —                |
| node_pools               | list(obj) | Node pool configurations     | 1× g6-standard-2 |
| additional_inbound_rules | list(obj) | Extra firewall rules         | []               |
| dns_zone_id              | number    | DNS zone ID (null=none)      | null             |
| dns_hostname             | string    | A record hostname            | ""               |
| tags                     | list      | Resource tags                | ["infra","k8s"]  |

Table C.3: k8s-platform variables

## C.4 shared-services Variables

| Variable             | Type   | Description                      | Default            |
|----------------------|--------|----------------------------------|--------------------|
| linode_token         | string | Linode API token (sensitive)     | —                  |
| region               | string | Linode region                    | eu-central         |
| gitlab_label         | string | GitLab CE Linode label           | porr-gitlab        |
| gitlab_instance_type | string | Linode plan for GitLab CE        | g6-standard-4      |
| vault_label          | string | Vault Linode label               | porr-vault         |
| vault_instance_type  | string | Linode plan for Vault            | g6-standard-2      |
| image                | string | Boot image                       | linode/debian12    |
| root_pass            | string | Root password (sensitive)        | —                  |
| ssh_keys             | list   | Authorized SSH keys              | []                 |
| vpc_id               | number | VPC ID from network-core         | —                  |
| subnet_id            | number | Subnet ID from network-core      | —                  |
| allowed_admin_ip     | list   | IPv4 CIDRs for SSH & Vault admin | ["0.0.0.0/0"]      |
| tags                 | list   | Resource tags                    | ["infra","shared"] |

| Variable | Type | Description | Default |
|----------|------|-------------|---------|
|----------|------|-------------|---------|

Table C.4: shared-services variables

## C.5 observability Variables

| Variable           | Type   | Description                   | Default          |
|--------------------|--------|-------------------------------|------------------|
| linode_token       | string | Linode API token (sensitive)  | —                |
| region             | string | Linode region                 | eu-central       |
| obs_label          | string | Linode instance label         | porr-obs         |
| instance_type      | string | Linode plan                   | g6-standard-2    |
| image              | string | Boot image                    | linode/debian12  |
| root_pass          | string | Root password (sensitive)     | —                |
| ssh_keys           | list   | Authorized SSH keys           | []               |
| vpc_id             | number | VPC ID from network-core      | —                |
| subnet_id          | number | Subnet ID from network-core   | —                |
| volume_size_gb     | number | Block storage size in GiB     | 200              |
| grafana_port       | number | Grafana HTTP port             | 3000             |
| allowed_grafana_ip | list   | IPv4 CIDRs for Grafana access | ["0.0.0.0/0"]    |
| tags               | list   | Resource tags                 | ["infra", "obs"] |

Table C.5: observability variables

# Appendix D

## HTTP Route Reference

| Method | Path                              | Description               | Role       |
|--------|-----------------------------------|---------------------------|------------|
| GET    | /                                 | Home / dashboard          | auth       |
| GET    | /catalog                          | Product catalog           | auth       |
| GET    | /catalog/{id}                     | Product detail            | auth       |
| GET    | /orders                           | My orders list            | auth       |
| GET    | /orders/new                       | New order form            | auth       |
| POST   | /orders/new                       | Submit order              | auth       |
| GET    | /orders/{id}                      | Order detail / status     | auth       |
| GET    | /projects                         | Projects list             | auth       |
| GET    | /infrastructure                   | Infrastructure overview   | auth       |
| POST   | /infrastructure/{id}/decommission | Trigger decommission      | auth       |
| GET    | /settings/profile                 | Profile / change password | auth       |
| GET    | /approvals                        | Approval queue            | admin      |
| POST   | /approvals/{id}/approve           | Approve order             | admin      |
| POST   | /approvals/{id}/reject            | Reject order              | admin      |
| GET    | /audit                            | Audit log                 | admin      |
| GET    | /audit/export                     | Export audit log          | admin      |
| GET    | /admin                            | Admin dashboard           | shop_admin |
| GET    | /admin/categories                 | Category list             | shop_admin |
| GET    | /admin/products                   | Product list              | shop_admin |
| GET    | /admin/products/new               | New product form          | shop_admin |
| GET    | /admin/products/{id}              | Edit product              | shop_admin |
| GET    | /admin/environments               | Environment list          | shop_admin |
| GET    | /admin/sources                    | GitLab sources list       | shop_admin |
| GET    | /admin/costcenters                | Cost center list          | shop_admin |
| GET    | /admin/users                      | User list                 | shop_admin |
| GET    | /admin/currencies                 | Currency management       | shop_admin |
| GET    | /admin/parameters                 | Global parameters         | shop_admin |
| GET    | /admin/smtp                       | SMTP configuration        | shop_admin |
| GET    | /admin/ai-config                  | AI translation config     | shop_admin |
| GET    | /admin/branding                   | Shop branding             | shop_admin |

| Method | Path           | Description   | Role   |
|--------|----------------|---------------|--------|
| GET    | /login         | Login page    | public |
| GET    | /auth/callback | OIDC callback | public |
| POST   | /logout        | Log out       | auth   |

Table D.1: HTTP route reference (selected)

# Appendix E

## Requirements Traceability Matrix

The following tables map each functional and non-functional requirement to the implementation location and the handbook section that describes it.

| ID    | Requirement                                 | Implementation                                                                              | Section              |
|-------|---------------------------------------------|---------------------------------------------------------------------------------------------|----------------------|
| FA-01 | User authentication (local + SSO)           | <a href="#">internal/service/impl/user.go</a> ,<br><a href="#">internal/handler/auth.go</a> | <a href="#">§8.8</a> |
| FA-02 | Role-based access control (admin, user, pm) | <a href="#">internal/middleware/auth.go</a> ,<br><a href="#">internal/model/user.go</a>     | <a href="#">§8.8</a> |
| FA-03 | Product catalog with categories             | <a href="#">internal/handler/catalog.go</a>                                                 | <a href="#">§4.3</a> |
| FA-04 | Configurable product parameters             | <a href="#">internal/model/parameter.go</a> ,<br>admin UI                                   | <a href="#">§4.6</a> |
| FA-05 | Shopping-cart order flow                    | <a href="#">internal/handler/order.go</a>                                                   | <a href="#">§13</a>  |
| FA-06 | Order approval workflow                     | <a href="#">internal/handler/approval.go</a>                                                | <a href="#">§13</a>  |
| FA-07 | GitLab webhook trigger                      | <a href="#">internal/service/impl/pipeline.go</a>                                           | <a href="#">§4.1</a> |
| FA-08 | Pipeline status polling                     | <a href="#">internal/polling/</a> ,<br><a href="#">internal/service/impl/pipeline.go</a>    | <a href="#">§4.5</a> |
| FA-09 | Infrastructure overview                     | <a href="#">internal/handler/infrastructure.go</a>                                          | <a href="#">§12</a>  |
| FA-10 | Decommission (destroy) flow                 | <a href="#">internal/handler/infrastructure.go</a>                                          | <a href="#">§15</a>  |
| FA-11 | Audit log                                   | <a href="#">internal/audit/</a> ,<br><a href="#">internal/handler/audit.go</a>              | <a href="#">§9</a>   |
| FA-12 | AI-powered product translation              | <a href="#">internal/service/impl/translation.go</a>                                        | <a href="#">§14</a>  |
| FA-13 | Multi-currency support                      | <a href="#">internal/service/impl/currency.go</a>                                           | <a href="#">§16</a>  |
| FA-14 | Branding configuration                      | <a href="#">internal/handler/admin.go</a> ,<br><a href="#">internal/model/branding.go</a>   | <a href="#">§7.5</a> |

| ID       | Requirement                          | Implementation                                                                 | Section |
|----------|--------------------------------------|--------------------------------------------------------------------------------|---------|
| NFA-01.1 | Docker deployment                    | <a href="#">Dockerfile</a> ,                                                   | §6.3    |
| NFA-01.2 | Kubernetes deployment                | <a href="#">deploy/docker-host/</a>                                            | §6.4    |
| NFA-01.3 | Single binary<br>(embed.FS)          | <a href="#">deploy/k8s/</a> manifests<br><a href="#">cmd/server/main.go</a> ,  | §3      |
| NFA-01.4 | Database migrations at<br>startup    | <a href="#">//go:embed</a> directives<br><a href="#">cmd/migrate/main.go</a> , | §11.5   |
| NFA-02.1 | HTTPS enforcement                    | <a href="#">internal/migrations/</a>                                           |         |
| NFA-02.2 | Secure session cookies               | Nginx config (reverse proxy,<br>TLS termination)                               | §6.3    |
| NFA-02.3 | bcrypt password<br>hashing           | <a href="#">internal/session/</a> ,<br>AES-GCM encrypted<br>HttpOnly cookies   | §3      |
| NFA-02.4 | Role-enforced routes                 | <a href="#">internal/service/impl/uses3go</a>                                  | §3      |
| NFA-03   | 25-language i18n                     | <a href="#">internal/middleware/auth.88</a>                                    | §8      |
| NFA-04   | Audit trail for all state<br>changes | <a href="#">internal/i18n/i18n.go</a>                                          | §A      |
|          |                                      | <a href="#">internal/audit/service.go</a>                                      | §9      |

Table E.1: Requirements traceability matrix

# Appendix F

## Compiling This Handbook

The handbook source is at [docs/handbook.tex](#). Compile with:

```
# From the repo root
make docs

# Or manually (run twice for ToC/references)
cd docs
pdflatex handbook.tex
pdflatex handbook.tex
```

**Required LaTeX packages:** lmodern, geometry, fancyhdr, xcolor, tikz, pgf, booktabs, longtable, tabularx, listings, tcolorbox, hyperref, enumitem, float, caption, graphicx, microtype, parskip, pdfscape, pifont.

All packages are available in TeX Live 2023+ and MiKTeX 23+.